

Natural Language Programming Project

Sentiment Analysis

Code Explanation

Imported Libraries :

```
# Importing Libraries Will Be Needed
import pandas as pd
import numpy as np
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer, WordNetLemmatizer
from nltk.tokenize import word_tokenize, sent_tokenize
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

#For embedding
from sentence_transformers import SentenceTransformer
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

# For Model Building
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.layers import LSTM, Bidirectional, Dropout, Dense, Embedding

from sklearn.model_selection import train_test_split

# Import metrics for detailed evaluation
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

- **Data Handling:** pandas, numpy for cleaning and processing data.
- **Text Preprocessing:** nltk, re for cleaning and tokenizing text.
- **Feature Extraction:** CountVectorizer, TfidfVectorizer for text to numeric conversion.
- **Embeddings:** SentenceTransformer, Tokenizer, pad_sequences for embedding creation.
- **Model Building:** Sequential, LSTM, Dense, Adam for neural networks.

- **Data Splitting:** `train_test_split` to divide training/testing sets.
- **Evaluation:** accuracy, precision, recall, f1 for performance metrics.

```
# Download NLTK resources
nltk.download('stopwords')
nltk.download('punkt') # For tokenization
nltk.download('wordnet') # For lemmatization
nltk.download('omw-1.4') # For WordNet Lemmatizer's language support

[2]

... [nltk_data] Downloading package stopwords to
[nltk_data]   C:\Users\elora\AppData\Roaming\nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data]   C:\Users\elora\AppData\Roaming\nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data]   C:\Users\elora\AppData\Roaming\nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
[nltk_data] Downloading package omw-1.4 to
[nltk_data]   C:\Users\elora\AppData\Roaming\nltk_data...
[nltk_data]   Package omw-1.4 is already up-to-date!

... True
```

Download NLTK Resources:

- stopwords: For removing common irrelevant words.
- punkt: For tokenizing text into words or sentences.
- wordnet: For lemmatization tasks.
- omw-1.4: Provides language support for WordNetLemmatizer.

```
# Load the dataset
data = pd.read_csv('Book1.csv')

# Display the first few rows of the dataset
print(data.head())
```

[3]

```
...               review sentiment
0  One of the other reviewers has mentioned that ...  positive
1  A wonderful little production. <br /><br />The...  positive
2  I thought this was a wonderful way to spend ti...  positive
3  Basically there's a family where a little boy ...  negative
4  Petter Mattei's "Love in the Time of Money" is...  positive
```

- **Load Dataset:** Reads a CSV file (`Book1.csv`) containing text reviews and sentiment labels.
- **Preview Data:** Displays the first few rows to confirm the structure:
 - **Columns:**
 - `review`: Text of the review.
 - `sentiment`: Sentiment label (positive/negative).

```
data['review'].duplicated().sum()
```

[4]

```
... np.int64(0)
```

- **Check for Duplicate Reviews:** Counts the number of duplicate entries in the `review` column.
 - **Result:** There are no duplicate reviews (0).
-

```
# Find duplicates in review column
duplicates = data[data['review'].duplicated()]

duplicates
```

[5]

```
...   review sentiment
```

- **Find Duplicate Reviews:** Identifies rows in the `review` column that are duplicates.
 - **Result:** No duplicates are found in the dataset (empty output).
-

```
data.isnull().sum()

[6]
... review      0
    sentiment    0
    dtype: int64
```

- **Check for Missing Values:** Verifies if there are any null or missing entries in the dataset.

```
data.describe()

[7]
...
      review  sentiment
count      1999      1999
unique      1999         2
top  I loved this movie! It was all I could do not ...  positive
freq              1      1005
```

- **Dataset Summary:** Provides statistical information about the data:
 - **Count:** 1,999 rows in both `review` and `sentiment` columns.
 - **Unique:** 1,999 unique reviews, indicating no duplicates.
 - **Top:** Most frequent review is "I loved this movie! It was all I could do not ..." and most frequent sentiment is "positive."
 - **Frequency:** Sentiment "positive" appears 1,005 times.

```
data['sentiment'].value_counts()

[8]
... sentiment
    positive    1005
    negative     994
    Name: count, dtype: int64
```

- **Sentiment Distribution:**
 - **Positive:** 1,005 reviews.

- **Negative:** 994 reviews.
 - The dataset is nearly balanced between the two sentiment classes.
-

```
✓ Data Cleaning and Preprocessing

# Removing duplicates in review column
data = data.drop_duplicates(subset=['review'])

[9]

data['review'].duplicated().sum()

[10]

... np.int64(0)
```

- **Remove Duplicates:** Eliminates duplicate entries in the `review` column using `drop_duplicates`.
- **Verify Duplicates:** Confirms no duplicate reviews remain in the dataset (0 duplicates).
- Ensures data integrity for analysis.

```

def clean_text(text):
    # Lowercase
    text = text.lower()

    # Remove HTML tags
    text = re.sub(r'<[>]+>', '', text)

    # Remove URLs
    text = re.sub(r'https?://\S+', '', text)

    # Remove special characters (punctuations) and numbers
    text = re.sub(r"[^a-zA-Z\s]", ' ', text)

    # Remove stopwords
    stop_words = set(stopwords.words('english'))
    text = ' '.join(word for word in text.split() if word not in stop_words)

    # Tokenization (word-level)
    word_tokens = word_tokenize(text) # Split into words

    # Stemming
    stemmer = PorterStemmer()
    stemmed_words = [stemmer.stem(word) for word in word_tokens]

    # Lemmatization
    lemmatizer = WordNetLemmatizer()
    lemmatized_words = [lemmatizer.lemmatize(word) for word in stemmed_words]

    # Combine the final processed words into a single string
    final_text = ' '.join(lemmatized_words)
    return final_text

# Apply the cleaning function to the review column
data['cleaned_review'] = data['review'].apply(clean_text)

```

[11]

- **Text Cleaning Function:**

- Converts text to lowercase for consistency.
- Removes HTML tags, URLs, special characters, and numbers for clarity.
- Removes stopwords (e.g., "the," "and") to focus on meaningful words.
- Tokenizes text into individual words.
- Applies stemming (reduces words to root form) and lemmatization (converts words to base form).
- Combines processed words back into a single string.

- **Application:** Cleans all reviews in the `review` column and stores the results in a new column called `cleaned_review`.

```

print(data.head())

[12]

...
      review sentiment \
0  One of the other reviewers has mentioned that ... positive
1  A wonderful little production. <br /><br />The... positive
2  I thought this was a wonderful way to spend ti... positive
3  Basically there's a family where a little boy ... negative
4  Petter Mattei's "Love in the Time of Money" is... positive

      cleaned_review
0  one review mention watch oz episod hook right ...
1  wonder littl product film techniqu unassum old...
2  thought wonder way spend time hot summer weeke...
3  basic famili littl boy jake think zombi closet...
4  petter mattei love time money visual stun film...

```

- **Data Preview:**

- Displays the first few rows of the dataset after cleaning.
- `review`: Original text reviews with raw formatting (HTML tags, stopwords, etc.).
- `cleaned_review`: Preprocessed reviews, cleaned and ready for analysis.
- The `cleaned_review` column shows simplified and processed text for better model input.

Text Representation

```

reviews = data['cleaned_review']
sentiments = data['sentiment'] # Target labels

[ ]

# 1. Bag of Words (BoW)
bow_vectorizer = CountVectorizer()
bow_features = bow_vectorizer.fit_transform(reviews)

print("BoW Shape:", bow_features.shape)
print("Sample BoW Vector:", bow_features[0].toarray())

[ ]

... BoW Shape: (1999, 17155)
     Sample BoW Vector: [[0 0 0 ... 0 0 0]]

```

- **Text Representation:**

- Converts cleaned reviews into numerical features for model training.

- **Bag of Words (BoW):**

- Uses `CountVectorizer` to create a sparse matrix of word frequencies.
- **Shape:** (1999, 17155) - 1999 reviews and 17,155 unique words.
- **Sample Vector:** Represents the word frequency of a single review as an array.

- **Purpose:** Transforms text data into a format suitable for machine learning models.


```
# 2. TF-IDF
tfidf_vectorizer = TfidfVectorizer()
tfidf_features = tfidf_vectorizer.fit_transform(reviews)

print("TF-IDF Shape:", tfidf_features.shape)
print("Sample TF-IDF Vector:", tfidf_features[0].toarray())

[ ]

... TF-IDF Shape: (1999, 17155)
Sample TF-IDF Vector: [[0. 0. 0. ... 0. 0. 0.]]
```

- **TF-IDF Representation:**
 - Converts cleaned reviews into numerical features using the **TF-IDF (Term Frequency-Inverse Document Frequency)** technique.
 - **Details:**
 - **Shape:** (1999, 17155) - 1999 reviews and 17,155 unique words.
 - **Sample TF-IDF Vector:** Represents the importance of each word in a review as a weighted value.
 - **Purpose:** Captures the importance of words in reviews relative to the dataset, making it more informative than simple word counts.
-

```
# 3. N-Grams (bi-grams or tri-grams)
ngram_vectorizer = CountVectorizer(ngram_range=(2, 2)) # Bi-grams
ngram_features = ngram_vectorizer.fit_transform(reviews)

print("N-Gram Shape:", ngram_features.shape)
print("Sample N-Gram Vector:", ngram_features[0].toarray())

[ ]

... N-Gram Shape: (1999, 177377)
Sample N-Gram Vector: [[0 0 0 ... 0 0 0]]
```

- **N-Grams Representation:**
 - Uses `CountVectorizer` to create features based on bi-grams (pairs of consecutive words).
- **Details:**
 - **Shape:** (1999, 177377) - 1999 reviews and 177,377 unique bi-grams.
 - **Sample N-Gram Vector:** Represents the frequency of bi-grams in a single review.
- **Purpose:** Captures contextual word pairs, providing more insight into word relationships within reviews.


```

# Load pre-trained BERT-based SentenceTransformer model
model = SentenceTransformer('paraphrase-MiniLM-L6-v2')

# Generate embeddings for all reviews
data['embeddings'] = data['cleaned_review'].apply(lambda x: model.encode(x))

# Convert to NumPy array for use in models
review_embeddings = np.array(data['embeddings'].tolist())
print("Embeddings Shape:", review_embeddings.shape)

[17]
... Embeddings Shape: (1999, 384)

```

- **Sentence Embeddings:**
 - Uses the pre-trained SentenceTransformer model (paraphrase-MiniLM-L6-v2) to generate dense semantic embeddings for each cleaned review.
 - **Details:**
 - **Embedding Shape:** (1999, 384) - 1999 reviews represented as 384-dimensional vectors.
 - Converts textual data into meaningful numerical representations capturing semantic context.
 - **Purpose:** Provides a powerful and context-aware representation of reviews, suitable for advanced machine learning models.
-

Padding

```

# Step 1: Initialize the tokenizer
tokenizer = Tokenizer(num_words=5000)

# Step 2: Fit the tokenizer on the cleaned review data
tokenizer.fit_on_texts(data['cleaned_review'])

# Step 3: Convert text to sequences of integers
sequences = tokenizer.texts_to_sequences(data['cleaned_review'])

# Step 4: Define the maximum sequence length
max_length = 500

# Step 5: Apply padding to standardize sentence lengths
padded_sequences = pad_sequences(sequences, maxlen=max_length, padding='post', truncating='post')

[18]

```

- **Padding Process:**
 1. **Initialize Tokenizer:** Limits the vocabulary to the top 5,000 most frequent words.
 2. **Fit Tokenizer:** Maps words in the cleaned reviews to integer indices.
 3. **Convert to Sequences:** Transforms text reviews into sequences of integers.
 4. **Define Maximum Length:** Sets a fixed sequence length (500).
 5. **Apply Padding:** Adds zeros to shorter sequences to standardize all to the same length.
- **Purpose:** Ensures all input sequences have uniform length for compatibility with machine learning models.

Model Building

```
# Step 6: Define the model architecture
vocab_size = len(tokenizer.word_index) + 1 # Vocabulary size from tokenizer
model = Sequential([
    Embedding(input_dim=vocab_size, output_dim=128, input_length=max_length), # Embedding layer
    Bidirectional(LSTM(64, return_sequences=False)), # Bidirectional LSTM
    Dropout(0.5), # Dropout for regularization
    Dense(1, activation='sigmoid') # Output layer for binary classification
])
```

- **Model Architecture:**
 - **Embedding Layer:** Converts words to dense vectors of size 128 based on the vocabulary size and input sequence length.
 - **Bidirectional LSTM:** Processes text in both forward and backward directions with 64 units, capturing context from both ends.
 - **Dropout:** Adds a 50% dropout rate to prevent overfitting.
 - **Dense Layer:** A single neuron with a sigmoid activation function for binary classification (outputting probabilities for positive or negative sentiment).
- **Purpose:** Builds a robust and context-aware model suitable for sentiment analysis.

Train the model

```
# Define X and y for training and testing
# X: Padded sequences (input features)
X = padded_sequences

# y: Sentiment labels (target)
# Convert sentiment into binary format: 1 for 'positive', 0 for 'negative'
y = (data['sentiment'] == 'positive').astype(int) # Ensure 'sentiment' column exists in the dataset

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=42)

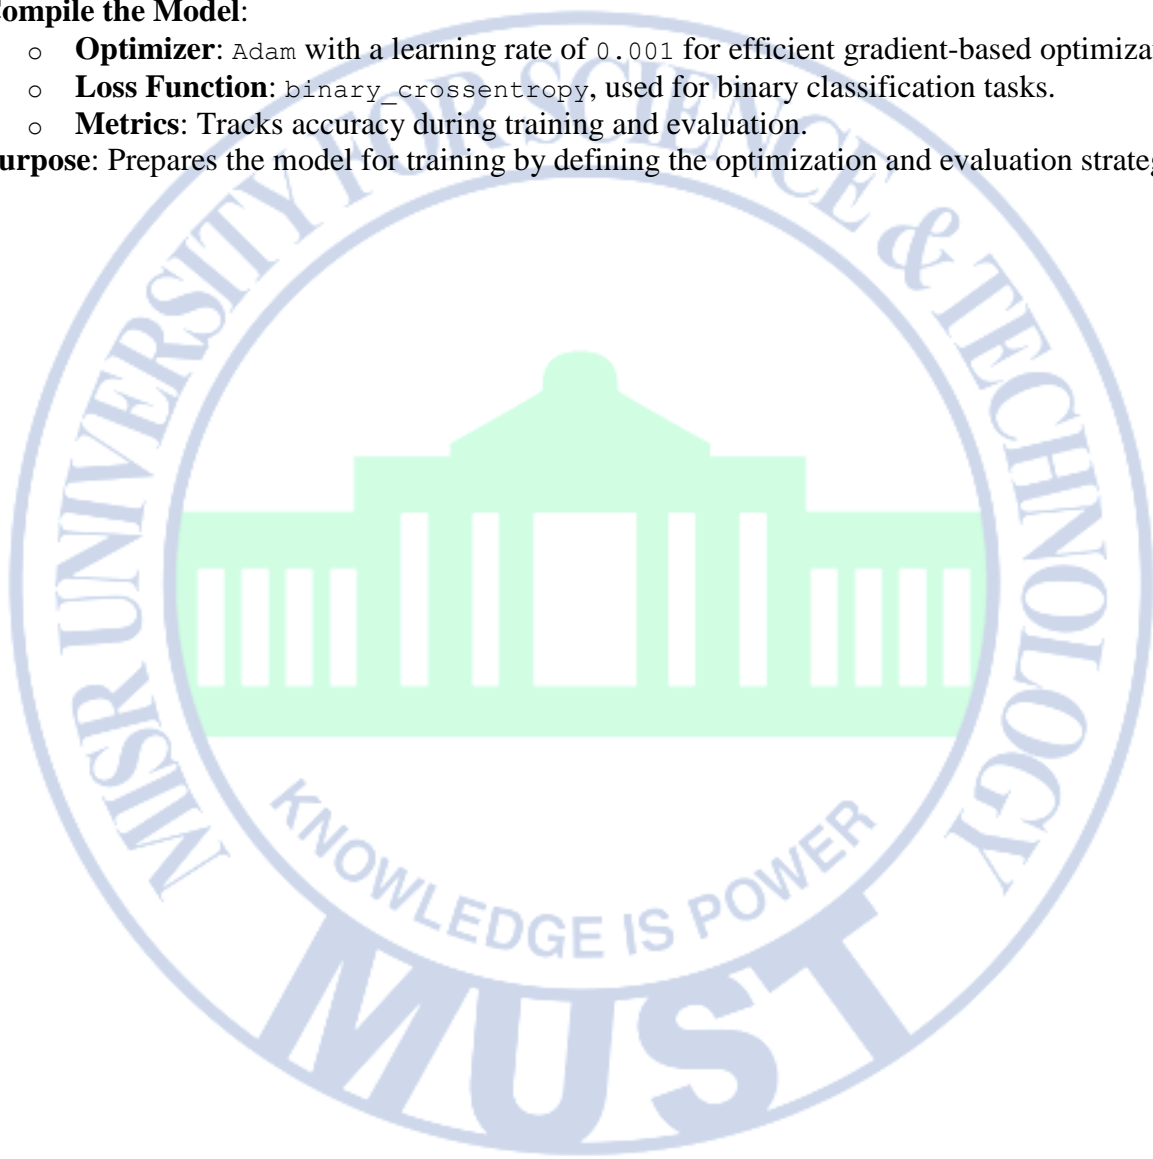
# Display shapes to confirm
print(f"Training Data Shape: X_train={X_train.shape}, y_train={y_train.shape}")
print(f"Testing Data Shape: X_test={X_test.shape}, y_test={y_test.shape}")
```

```
[27]
... Training Data Shape: X_train=(1199, 500), y_train=(1199,)
Testing Data Shape: X_test=(800, 500), y_test=(800,)
```

- **Train-Test Split:**
 - **Input Features (x):** Uses padded_sequences as input data.
 - **Target Labels (y):** Converts sentiment to binary format (1 for positive, 0 for negative).
 - Splits the data into:
 - **Training Set:** 60% of the data for training the model (x_train, y_train).
 - **Testing Set:** 40% of the data for evaluating the model (x_test, y_test).
- **Shapes:**
 - Training Data: 1,199 reviews with 500 features each.
 - Testing Data: 800 reviews with 500 features each.
- **Purpose:** Prepares the dataset for training and ensures separate evaluation data for unbiased model testing.

```
# Compile the model
model.compile(optimizer=Adam(learning_rate=0.001), # Optimizer
              loss='binary_crossentropy',         # Loss function for binary classification
              metrics=['accuracy'])               # Accuracy metric for evaluation
```

- **Compile the Model:**
 - **Optimizer:** Adam with a learning rate of 0.001 for efficient gradient-based optimization.
 - **Loss Function:** `binary_crossentropy`, used for binary classification tasks.
 - **Metrics:** Tracks accuracy during training and evaluation.
- **Purpose:** Prepares the model for training by defining the optimization and evaluation strategies.



MISR UNIVERSITY
FOR SCIENCE & TECHNOLOGY

```

from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau

# Train the model with early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)
lr_scheduler = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=2, verbose=1)

# Train the model with validation split, early stopping, and learning rate scheduler
history = model.fit(X_train, y_train,
                    epochs=15,           # Training data
                    batch_size=32,       # Number of epochs
                    validation_split=0.4, # Batch size
                    callbacks=[early_stopping, lr_scheduler], # Validation split (40% of training data)
                    verbose=1)           # Verbosity mode

# Display the model summary
print(model.summary())

```

[29]

```

... Epoch 1/15
23/23 ----- 14s 439ms/step - accuracy: 0.8906 - loss: 0.3097 - val_accuracy: 0.7729 - val_loss: 0.4833 - learning_rate: 0.0010
Epoch 2/15
23/23 ----- 10s 448ms/step - accuracy: 0.9618 - loss: 0.1581 - val_accuracy: 0.8021 - val_loss: 0.4788 - learning_rate: 0.0010
Epoch 3/15
23/23 ----- 10s 448ms/step - accuracy: 0.9868 - loss: 0.0922 - val_accuracy: 0.8062 - val_loss: 0.5250 - learning_rate: 0.0010
Epoch 4/15
23/23 ----- 0s 414ms/step - accuracy: 0.9967 - loss: 0.0386
Epoch 4: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.
23/23 ----- 11s 476ms/step - accuracy: 0.9966 - loss: 0.0386 - val_accuracy: 0.8083 - val_loss: 0.5945 - learning_rate: 0.0010
Epoch 5/15
23/23 ----- 10s 452ms/step - accuracy: 0.9985 - loss: 0.0221 - val_accuracy: 0.8146 - val_loss: 0.5929 - learning_rate: 2.0000e-04
...
Model: "sequential"
...

```

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 500, 128)	2,198,784
bidirectional (Bidirectional)	(None, 128)	98,816
dropout (Dropout)	(None, 128)	0
dense (Dense)	(None, 1)	129

```

... Total params: 6,893,189 (26.30 MB)
...
... Trainable params: 2,297,729 (8.77 MB)
...
... Non-trainable params: 0 (0.00 B)
...
... Optimizer params: 4,595,460 (17.53 MB)
...
... None

```

- **Training:**
 - EarlyStopping: Stops training if validation loss does not improve for 3 epochs and restores the best weights.
 - ReduceLROnPlateau: Reduces learning rate after 2 epochs of no improvement in validation loss.
 - Model trains for up to 15 epochs with a batch size of 32, but stops early due to callbacks.
- **Training Metrics:**
 - Accuracy improves steadily, reaching 99.85% on training data.
 - Validation accuracy reaches 81.46% after early stopping.
 - Learning rate is reduced during training to fine-tune performance.
- **Model Summary:**
 - **Embedding Layer:** Maps input into a dense vector space with 2,198,784 parameters.
 - **Bidirectional LSTM:** Processes sequences in both forward and backward directions with 98,816 parameters.
 - **Dropout Layer:** Adds no parameters; used for regularization.
 - **Dense Output Layer:** Final layer with 1 neuron for binary classification.
- **Model Details:**
 - **Total Parameters:** 6,893,189.
 - **Trainable Parameters:** 2,297,729.
 - **Optimizer Parameters:** 4,595,460.
- **Purpose:** Achieves high accuracy while preventing overfitting through callbacks and regularization.

Evaluate the model

```
# Evaluate on test data
loss, accuracy = model.evaluate(X_test, y_test, verbose=0)

# Predict on test data
y_pred = model.predict(X_test)
y_pred_classes = (y_pred > 0.5).astype(int) # Convert probabilities to binary labels

# Calculate metrics
accuracy = accuracy_score(y_test, y_pred_classes)
precision = precision_score(y_test, y_pred_classes)
recall = recall_score(y_test, y_pred_classes)
f1 = f1_score(y_test, y_pred_classes)

# Display metrics
print(f"Test Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1-Score: {f1:.2f}")
```

[38]

... 25/25 ————— 2s 83ms/step
Test Accuracy: 0.84
Precision: 0.84
Recall: 0.85
F1-Score: 0.85

- **Model Evaluation:**
 - Evaluates the trained model on test data to calculate **loss** and **accuracy**.
 - Test accuracy: 84%.
- **Predictions:**
 - Converts predicted probabilities to binary labels (0 or 1) using a threshold of 0.5.
- **Performance Metrics:**
 - **Accuracy:** 84% - Overall correctness of predictions.
 - **Precision:** 84% - Proportion of true positives among predicted positives.
 - **Recall:** 85% - Proportion of true positives identified correctly.
 - **F1-Score:** 85% - Balance between precision and recall.
- **Purpose:** Confirms the model's effectiveness in classifying sentiments with good accuracy and balanced precision/recall.